

trunSVD_xgboost_shap_windows

June 11, 2023

```
[ ]: import warnings

warnings.filterwarnings("ignore",category=DeprecationWarning)

[1]: # Import necessary libraries
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import make_column_transformer
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
import shap

[2]: # Load a local dataset
# Replace 'path_to_your_dataset' with the actual path to your local dataset
df = pd.read_excel(r'C:\Users\huangchuhuan\Downloads\    .xlsx', header = 0,
↳skiprows = [0])
# focus on
df = df.dropna(subset = [' '])
# ignore rows with nan
df = df.dropna(axis=0)
# convert % symbol to real values
for column in df.columns:
    if df[column].dtype == np.object:
        if df[column].str.contains('%').any():
            df[column] = df[column].str.rstrip('%').astype(float)*100.0/(100.
↳0**2)

`np.object` is a deprecated alias for the builtin `object`. To silence this
warning, use `object` by itself. Doing this will not modify any behavior and is
safe.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations

[3]: # Select a random sample of 100 rows
df_sample = df.sample(n=100, random_state=1)
```

```
df_sample.head()
```

				\n()	\
4413	040016.OF			2010-05-11		2018-11-12
2098	008208.OF			2019-12-19		2019-12-19
4659	163807.OF		A	2009-04-03		2015-05-28
1497	005260.OF	, ,	A	2017-12-15		2017-12-15
4435	070002.OF			2003-07-09		2018-12-04

Wind	()4	.11	\
4413	6.315900	0.0	...	Beta	Beta
2098	12.784945	0.0	...	Beta	Beta
4659	34.999295	0.0	...	Beta	Beta
1497	0.321281	0.0	...	Beta	Beta
4435	30.490955	0.0	...	Beta	Beta

4413	,	,	0
2098	,	,	0
4659	,	,	0
1497	,	,	0
4435	,	,	

```
[5 rows x 83 columns]
```

```
# Autonomously choose a numerical variable as the target
# Here we choose 'ROE_TTM%' as an example
y = df_sample[['ROE_TTM%']]
# Drop the target variable from the dataframe
df_sample = df_sample.drop('ROE_TTM%', axis=1)
```

```
# Get lists of numerical and categorical columns
num_cols = [col for col in df_sample.columns if df_sample[col].dtype in_
↳ ['int64', 'float64']]
cat_cols = [col for col in df_sample.columns if df_sample[col].dtype ==_
↳ 'object']
```

```
# Convert all categorical columns to string type
df_sample[cat_cols] = df_sample[cat_cols].astype(str)

# Now define the preprocessor
preprocessor = make_column_transformer(
    (StandardScaler(), num_cols), # standardize numerical features
    (OneHotEncoder(handle_unknown='ignore', sparse=False), cat_cols) # one-hot
    ↪ encode categorical features
)
```

```
# Preprocess the data
X = preprocessor.fit_transform(df_sample)
```

```
[8]: from sklearn.decomposition import TruncatedSVD
```

```
[9]: num_cols = 16
     num_rows = 100
```

```
[10]: data = [[np.nan]*num_cols for _ in range(num_rows)]
      train_score_df = pd.DataFrame(data)
      test_score_df = pd.DataFrame(data)
```

```
[11]: for i in range(4,num_cols+4):
      for j in range(num_rows):
          # Apply TruncatedSVD
          svd = TruncatedSVD(n_components=i) # specify the number of components,
          ↪can be adjusted based on your needs
          X_svd = svd.fit_transform(X)
          # Split the data into training and test sets
          X_train, X_test, y_train, y_test = train_test_split(X_svd, y,
          ↪test_size=0.2, random_state=42)

          # Train an XGBoost regressor
          xgb = XGBRegressor(objective='reg:squarederror', random_state=42)
          xgb.fit(X_train, y_train)

          # Evaluate the model
          train_score_df.iat[j,i-4] = xgb.score(X_train, y_train)
          test_score_df.iat[j,i-4] = xgb.score(X_test, y_test)
```

```
[12]: test_score_df.head(10)
```

```
[12]:
```

	0	1	2	3	4	5	6 \
0	0.399332	0.707178	0.731261	0.683504	0.759852	0.656395	0.637740
1	0.401846	0.712542	0.755513	0.680714	0.757530	0.645657	0.588085
2	0.401846	0.707178	0.760670	0.671762	0.758432	0.666944	0.607315
3	0.399332	0.708568	0.757364	0.675555	0.759756	0.666395	0.634048
4	0.401846	0.709355	0.757363	0.663761	0.759608	0.635076	0.659257
5	0.401846	0.707178	0.756860	0.675865	0.756873	0.669821	0.521219
6	0.399332	0.707178	0.757364	0.687093	0.757052	0.651431	0.625772
7	0.401846	0.707178	0.755915	0.675381	0.757320	0.654603	0.588539
8	0.401846	0.707178	0.758674	0.683531	0.753802	0.668210	0.623493
9	0.401846	0.711161	0.755506	0.678314	0.757818	0.626702	0.640143

	7	8	9	10	11	12	13 \
0	0.678626	0.703874	0.669170	0.641379	0.491338	0.467539	0.476261
1	0.676284	0.706438	0.645756	0.667462	0.478638	0.467924	0.489849

2	0.707286	0.667478	0.742182	0.636922	0.474465	0.502289	0.460614
3	0.686651	0.659073	0.684626	0.655692	0.498198	0.462394	0.469451
4	0.651343	0.749167	0.646477	0.639052	0.478371	0.458417	0.469359
5	0.678963	0.638780	0.671207	0.639060	0.481431	0.460126	0.491408
6	0.687651	0.660770	0.649684	0.715158	0.460258	0.487854	0.493682
7	0.634687	0.663398	0.644066	0.668449	0.481339	0.454811	0.486459
8	0.730357	0.737300	0.681620	0.735028	0.494197	0.496407	0.470187
9	0.683182	0.731398	0.657197	0.653669	0.490023	0.453435	0.528558

	14	15
0	0.489494	0.680423
1	0.477648	0.676704
2	0.467889	0.653162
3	0.553708	0.675854
4	0.447270	0.684364
5	0.482344	0.628213
6	0.474138	0.683142
7	0.546508	0.555334
8	0.466828	0.541357
9	0.480619	0.693483

```
[13]: import matplotlib.pyplot as plt
```

```
[14]: col_names = {}
      for column in test_score_df.columns:
          print(type(column))
          col_names[column] = column+4
      test_score_df = test_score_df.rename(col_names, axis=1)
```

```
<class 'int'>
<class 'int'>
<class 'int'>
<class 'int'>
<class 'int'>
<class 'int'>
<class 'int'>
<class 'int'>
<class 'int'>
<class 'int'>
<class 'int'>
<class 'int'>
<class 'int'>
<class 'int'>
<class 'int'>
```

```
[15]: print(col_names)
```

```
{0: 4, 1: 5, 2: 6, 3: 7, 4: 8, 5: 9, 6: 10, 7: 11, 8: 12, 9: 13, 10: 14, 11: 15,
12: 16, 13: 17, 14: 18, 15: 19}
```

```
[16]: test_score_df.head()
```


```
[16]:
```

	4	5	6	7	8	9	10 \
0	0.399332	0.707178	0.731261	0.683504	0.759852	0.656395	0.637740
1	0.401846	0.712542	0.755513	0.680714	0.757530	0.645657	0.588085
2	0.401846	0.707178	0.760670	0.671762	0.758432	0.666944	0.607315
3	0.399332	0.708568	0.757364	0.675555	0.759756	0.666395	0.634048
4	0.401846	0.709355	0.757363	0.663761	0.759608	0.635076	0.659257

	11	12	13	14	15	16	17 \
0	0.678626	0.703874	0.669170	0.641379	0.491338	0.467539	0.476261
1	0.676284	0.706438	0.645756	0.667462	0.478638	0.467924	0.489849
2	0.707286	0.667478	0.742182	0.636922	0.474465	0.502289	0.460614
3	0.686651	0.659073	0.684626	0.655692	0.498198	0.462394	0.469451
4	0.651343	0.749167	0.646477	0.639052	0.478371	0.458417	0.469359

	18	19
0	0.489494	0.680423
1	0.477648	0.676704
2	0.467889	0.653162
3	0.553708	0.675854
4	0.447270	0.684364

```
[17]: test_score_df.boxplot()
plt.show()
```



```
[18]: def get_column_with_largest_mean(df):
    # Calculate the mean values of each column
    column_means = df.mean()

    # Get the index of the column with the largest mean value
    largest_mean_index = column_means.idxmax()

    # Check for ties by comparing variance values
    column_variances = df.var()
    tiebreaker_indices = []
    largest_variance = df.var()[largest_mean_index]

    # Find indices of columns with the same mean value as the largest mean
    ↪column
    for col in df.columns:
        if column_means[col] == column_means[largest_mean_index]:
            if df.var()[col] < largest_variance:
                largest_variance = df.var()[col]
                tiebreaker_indices = [col]
            elif df.var()[col] == largest_variance:
                tiebreaker_indices.append(col)

    # Check for tiebreaker indices
    if tiebreaker_indices:
        # Return the index of the column with the smallest variance and
        ↪smallest index
        return min(tiebreaker_indices)
    else:
        # Return the index of the column with the largest mean value
        return largest_mean_index
```

```
[19]: n = get_column_with_largest_mean(test_score_df)
print(n)
```

8

```
[20]: #Choosing based on rank of average

# Apply TruncatedSVD
svd = TruncatedSVD(n_components=n) # specify the number of components, can be
    ↪adjusted based on your needs
X_svd = svd.fit_transform(X)
# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X_svd, y, test_size=0.2,
    ↪random_state=42)
```

```

# Train an XGBoost regressor
xgb = XGBRegressor(objective='reg:squarederror', random_state=42)
xgb.fit(X_train, y_train)

# Evaluate the model
train_score = xgb.score(X_train, y_train)
test_score = xgb.score(X_test, y_test)

print(f"Train score: {train_score}")
print(f"Test score: {test_score}")

```

Train score: 0.9999999785500552

Test score: 0.7570926899283865

```
[21]: shap.initjs()
```

<IPython.core.display.HTML object>

```
[22]: explainer = shap.TreeExplainer(xgb)
```

```
[23]: shap_values = explainer.shap_values(X_train)
```

ntree_limit is deprecated, use `iteration_range` or model slicing instead.

```
[24]: shap.summary_plot(shap_values,X_train)
```

